

# Image Skeletonisation

A performance comparison with Thrust


Fall 2014

Marco Antognini



ÉCOLE POLYTECHNIQUE  
FÉDÉRALE DE LAUSANNE

# What is it?





A method to reduce  
a **binary image** to  
its **minimal representation**


Useful for

- blood analysis
- classification of fingerprints
- text recognition
- geometrical shape analysis
- ...

# What is it?



# Serial Algorithm



**H** a sub-iteration

- blue  $\Rightarrow$  first sub-cycle
- red  $\Rightarrow$  second sub-cycle

**I** the marked pixels are deleted

For each sub-cycle

For each pixel  $p$  in the image

- If some predicates on the 8 direct neighbours don't hold
- then  $p$  is **marked** for deletion

After  $M$  **iterations** the image has converged to its skeleton

$M$  depends on the shape

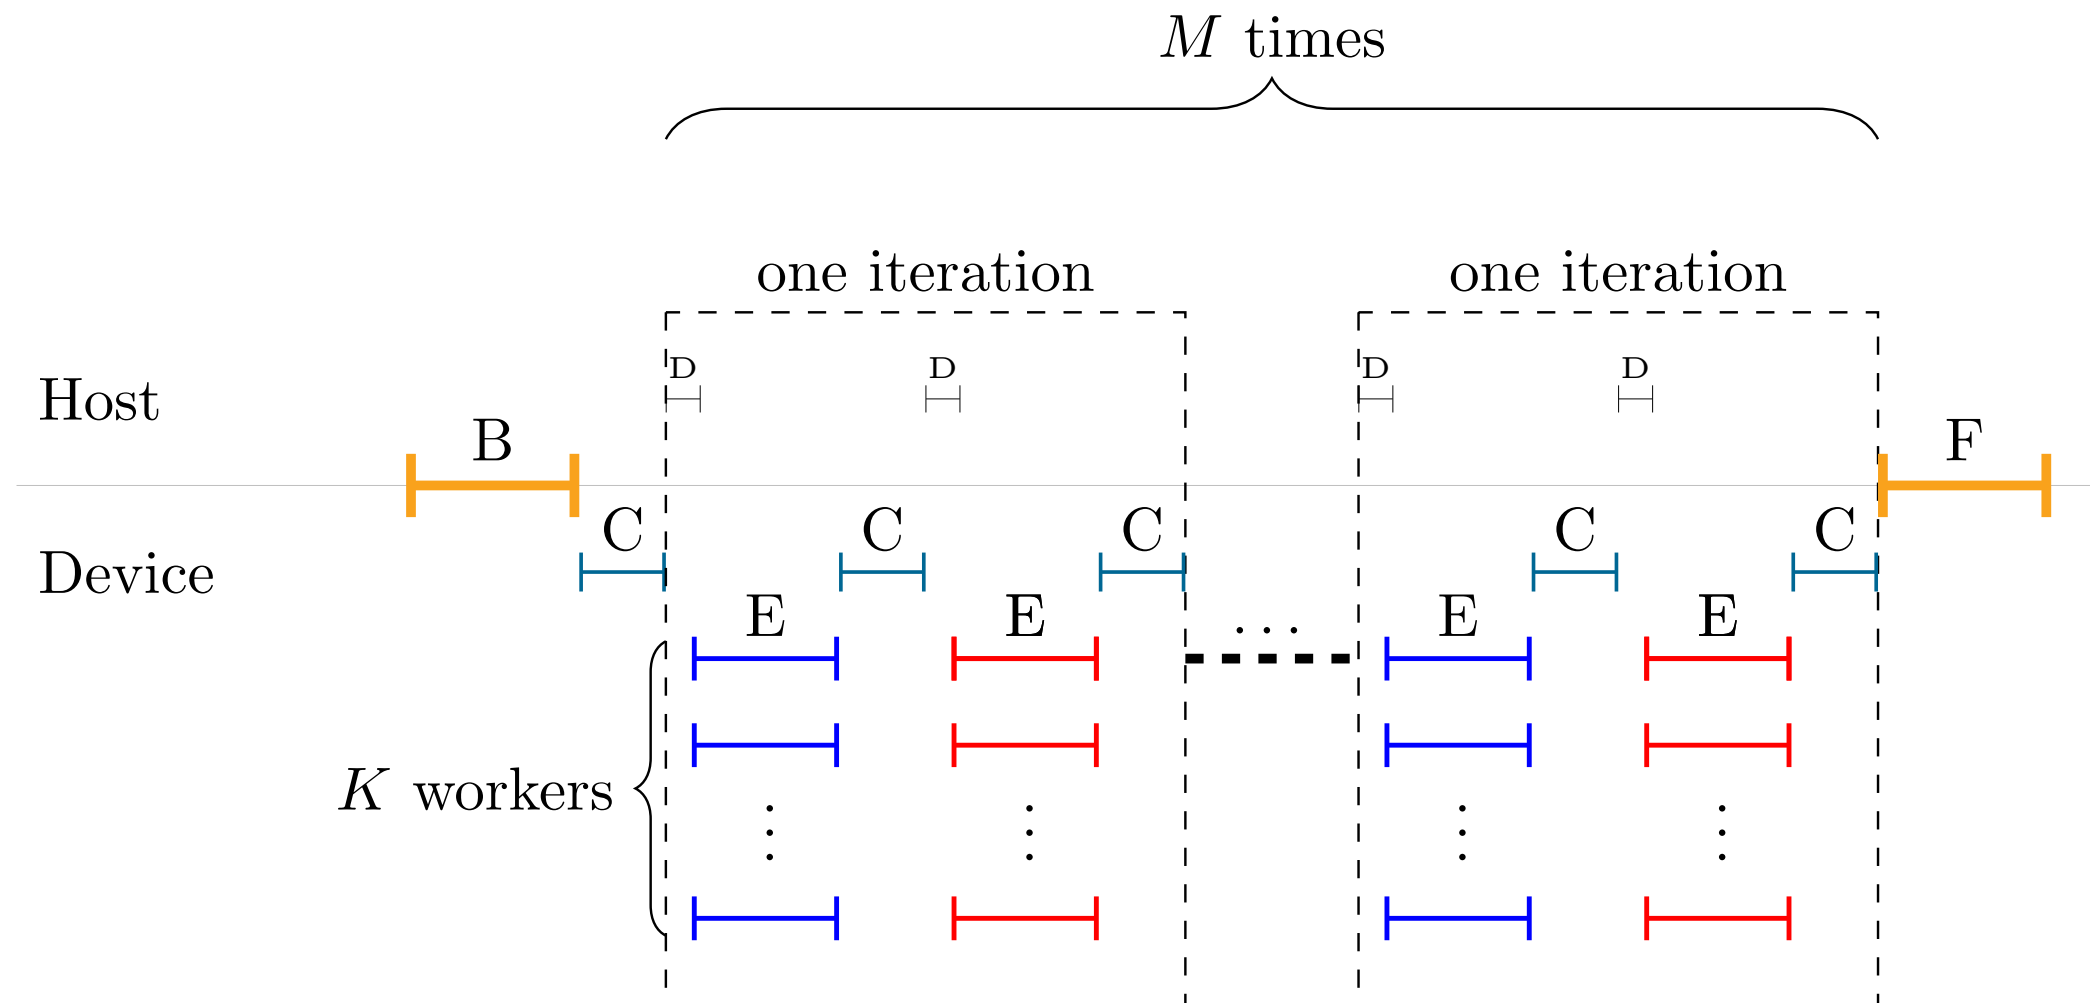
- the bigger the widest arm of the shape is, the bigger  $M$  will be

# Implementation: Thrust

- ➔ A C++03 general purpose library by Nvidia
- ➔ GPU and CPU backends **with the same code**
  - CUDA, OpenMP, Threading Building Blocks
- ➔ API very similar to STL
- ➔ Based on high level **functional algorithms**
  - *transform, reduce, copy, count, search, partition, ...*
- ➔ But can't do everything...
  - e.g. matrix multiplication



# Parallel Algorithm



**Each pixel can be marked regardless of the other pixels** in a sub-cycle

- ➔ We use  $K$  workers to split the job
- ➔ With double buffering (buffer ❶ and ❷)
- ➔ Instead of marking we delete in ❷

**B** send data to *device* memory, buffer ❶

**C** synchronise buffers ❶ and ❷

**D/E** launch kernel function on *device*

▸ blue  $\Rightarrow$  first sub-cycle, red  $\Rightarrow$  second sub-cycle

**F** copy data back to *host*

# Theoretical Speedup

## Work split on $K$ workers

- ➔ CUDA on GeForce GT 650M, Kepler GK110 architecture
  - 2 Multiprocessors (MP) @ 900 MHz with 192 cores each
  - Max 2048 threads per MP understand: resource available for max 2048 threads
  - But only 4 warps executing instructions simultaneously per MP
  - $\Rightarrow 2 \text{ MP} * 4 \text{ warps} * 32 \text{ threads} / \text{warps} = \mathbf{K = 256}$
- ➔ OMP/TBB on Core i7 3720QM
  - 4 cores @ 2600 MHz with Hyper-Threading Technology
  - $\Rightarrow \mathbf{K = 8}$

## Clock ratio

- $\Rightarrow \lambda \approx 2.88$

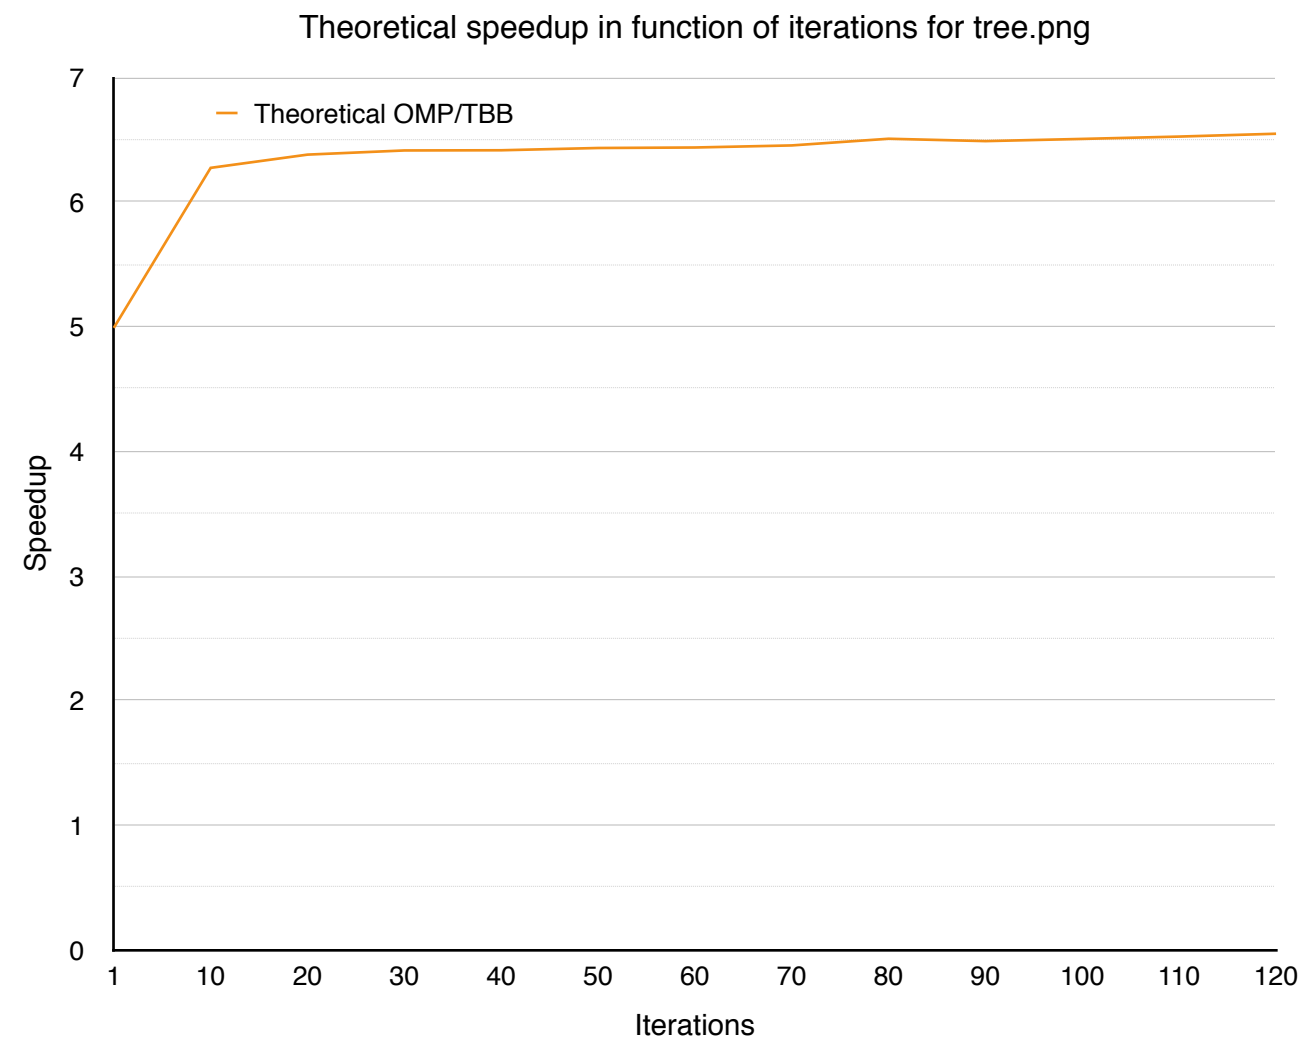
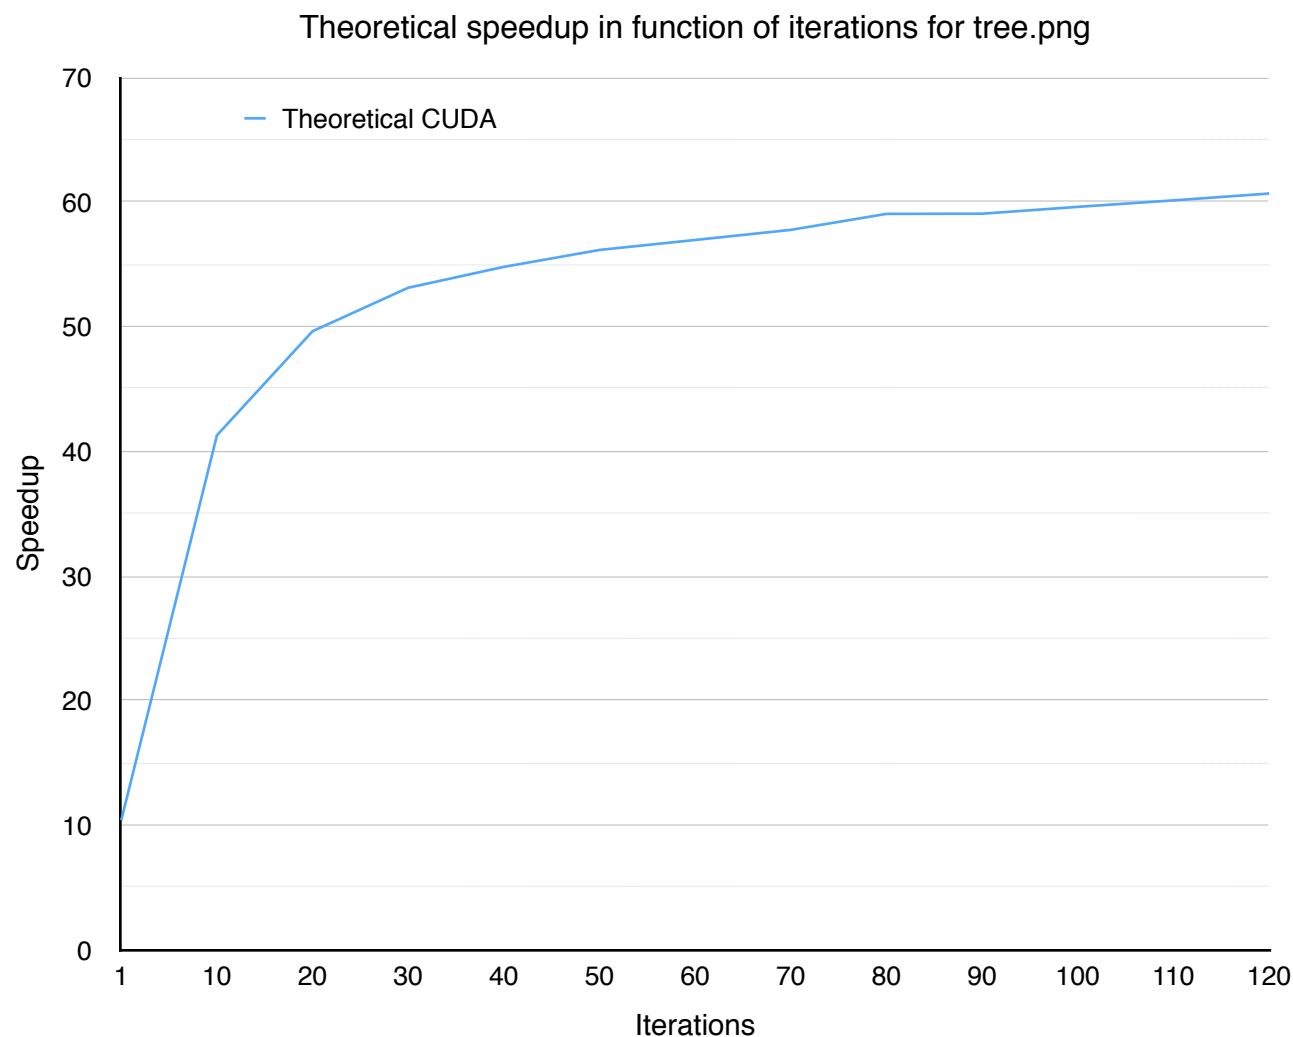
Memory speed depends on transferred size

- See final report

# Theoretical Speedup

$$Speedup(N, M, K) = \frac{t_{ser}}{t_{par}} = \frac{t_{ser}}{\frac{N}{H \rightarrow D_s} + \frac{N}{D \leftrightarrow D_s} + 2M \left( \frac{\lambda}{K} t_{sub} + \frac{N}{D \leftrightarrow D_s} \right) + \frac{N}{D \rightarrow H_s}}$$

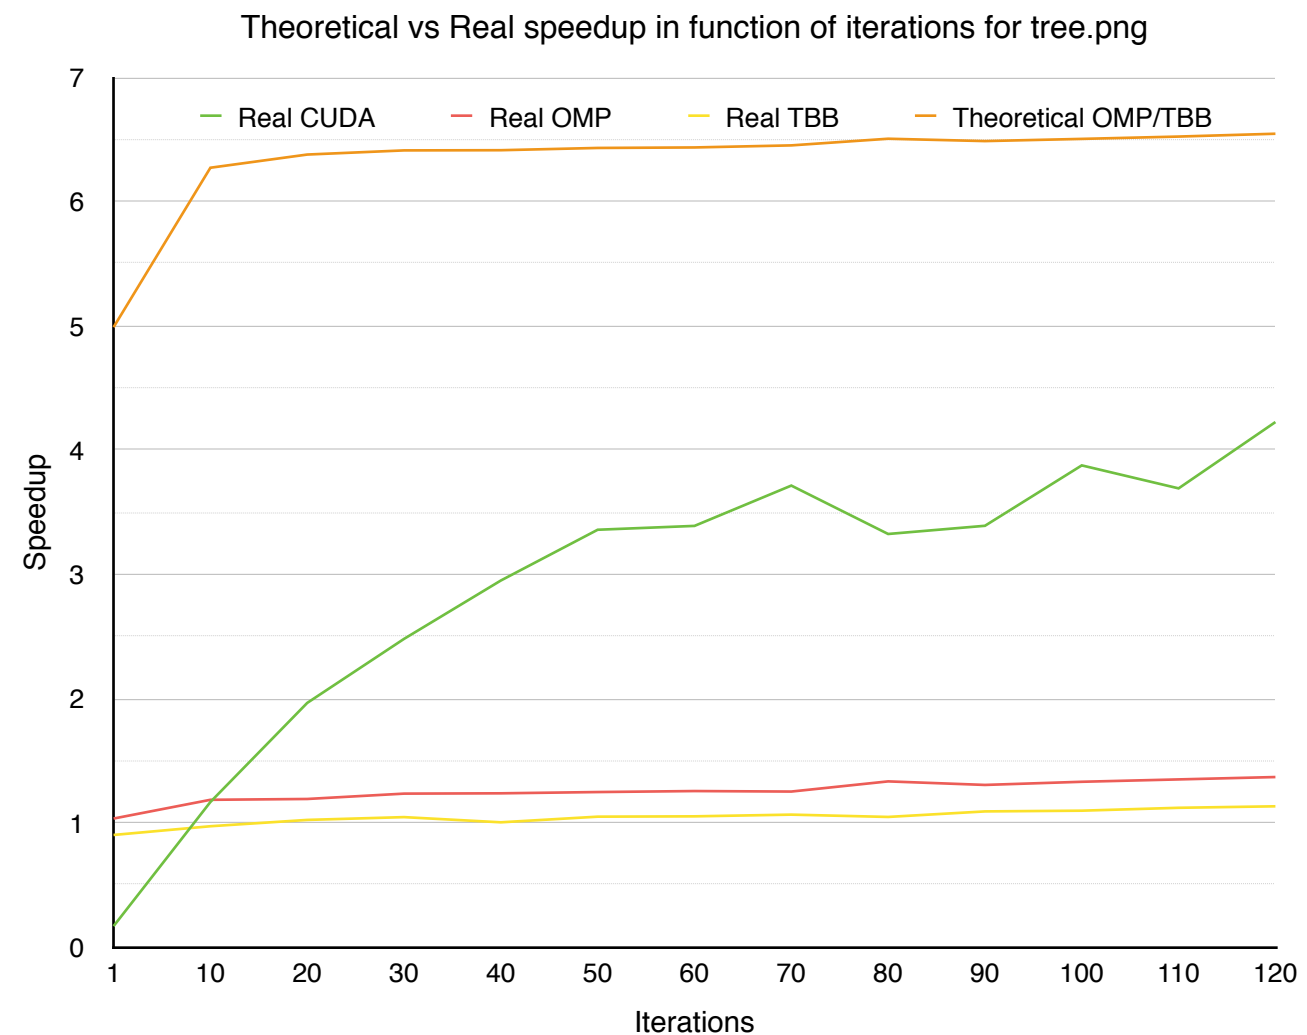
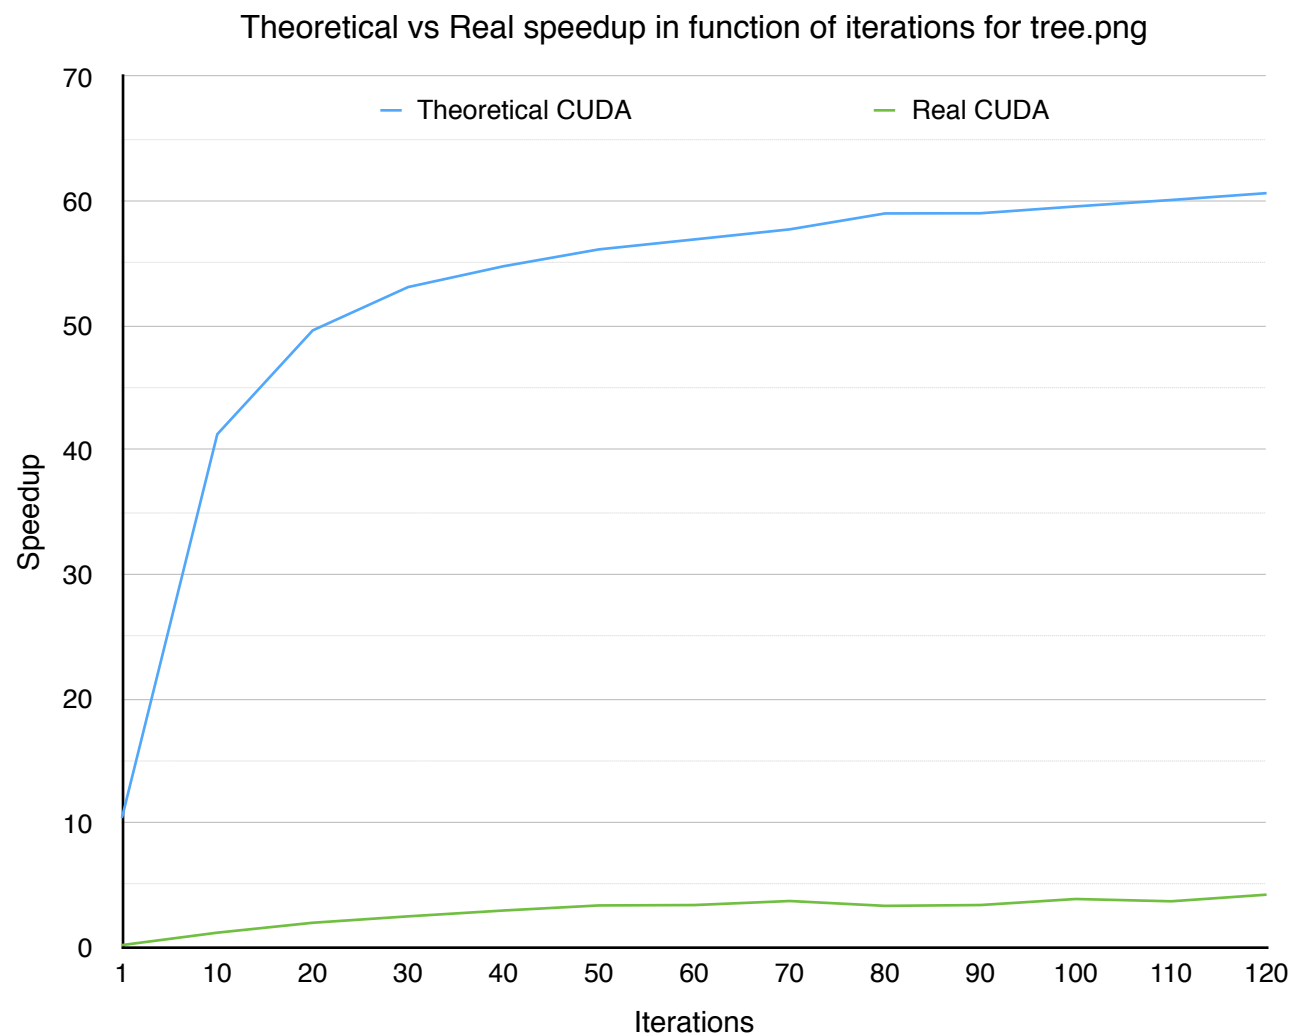
- The number of pixels,  $N$
- The number of required iterations,  $M$
- The number of worker threads,  $K$





# The Reality

- ➔ The speedups are **much smaller**
- ➔ CUDA still offer significant speedup, though
- ➔ CUDA can be slower when  $M$  is too small
  - Can't always compensate for transfer times and slower GPU clock



# The Reality: Why?

**The model doesn't take into account several important points**

- ➔ Resource availability
  - **Memory Banks**: queuing when memory bus is busy
  - **Cache Misses**: potentially huge delays
  - **Limited ALUs**: operations delayed when no available ALU
  - 32-bits VS 64-bits operations: up to 2/3 performance lost
- ➔ Thread divergence
- ➔ Data Structure  $\Leftrightarrow$  **Stridden Memory**
  - If not carefully designed, accessing data kill bandwidth
  - E.g. stride-two access imply a 50% bandwidth penalty

**Such model would be too difficult to use...**

# What's not there?

- ➔ Speedup VS Iterations
  - Other images show similar behaviour than *tree.png*
- ➔ Speedup VS Image Size
  - What happens when  $N$  varies but not  $M$ ?
- ➔ Memory Bandwidth
  - Speed varies with transferred size

Questions?